

# DSc Assignment 3

Vaibhav Mittal

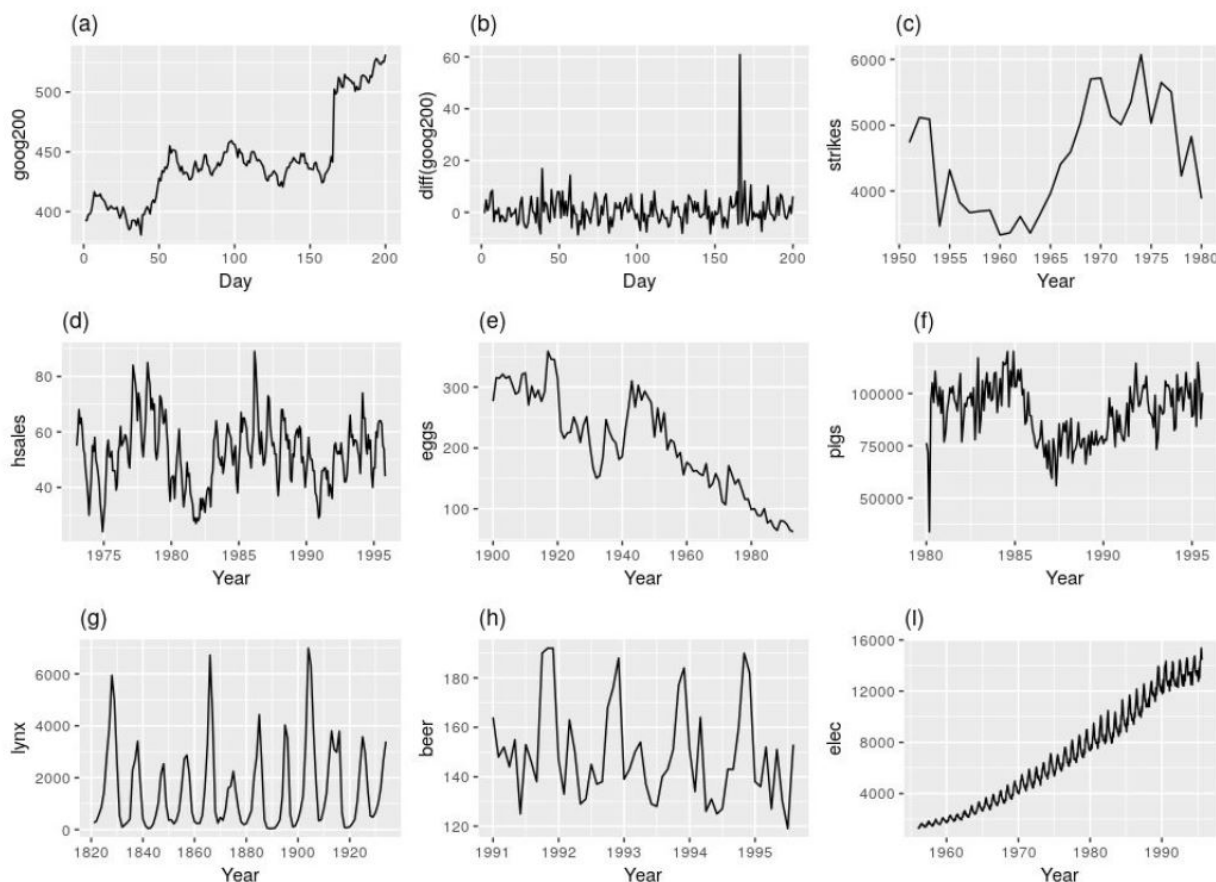
24 October 2019

## Time Series Analysis

Submitted by Vaibhav Mittal, MT18242.

### Question 1

```
setwd("E:/College/Sem3/DSc/Assign3/")
```



Question:

*a* cannot be considered stationary since the mean is changing. *b* can be considered stationary *iff* outliers aren't considered. Because of some instances of changing covariances, *b* isn't stationary. *c* isn't stationary.

*d* isn't stationary. It has some oscillatory behaviour too. *e* isn't stationary due to changing mean. *f* isn't stationary too due to the same reasons.

*g* seems stationary albeit it has some outliers. *h* is a bit less stationary since it has more outliers. There are some covariance changes too. *i* is not at all stationary.

## Question 2

```
library(quantmod)

## Loading required package: xts
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
##
## Registered S3 method overwritten by 'xts':
##   method      from
##   as.zoo.xts zoo
##
## Loading required package: TTR
##
## Registered S3 method overwritten by 'quantmod':
##   method      from
##   as.zoo.data.frame zoo
##
## Version 0.4-0 included new data defaults. See ?getSymbols.
getSymbols(c('AMZN', 'MSFT', 'GOOG'), src = 'yahoo', from = '2016-10-01')

## 'getSymbols' currently uses auto.assign=TRUE by default, but will
## use auto.assign=FALSE in 0.5-0. You will still be able to use
## 'loadSymbols' to automatically load data. getOption("getSymbols.env")
## and getOption("getSymbols.auto.assign") will still be checked for
## alternate defaults.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.
## [1] "AMZN" "MSFT" "GOOG"


We won't go into Bollinger Band charts and will only plot their adjusted time series.
plot(AMZN[,6])
```



```
plot(GOOG[,6])
```

GOOG[, 6]

2016-10-03 / 2019-10-25



```
plot(MSFT[,6])
```



All the stocks seem healthy in the fact that their means show an increasing trend.

```
amazon = c(AMZN[,6])
google = c(GOOG[,6])
microsoft = c(MSFT[,6])

amazon_1 = rep(0, length(amazon))
amazon_1[1:(length(amazon)-1)] = c(amazon[2:length(amazon)])
amazon_2 = rep(0, length(amazon))
amazon_2[1:(length(amazon)-2)] = c(amazon[3:length(amazon)])
amazon_3 = rep(0, length(amazon))
amazon_3[1:(length(amazon)-3)] = c(amazon[4:length(amazon)])
#plot(x_t, x_t_1)
cor(amazon, amazon_1)
```

```
##           [,1]
## AMZN.Adjusted 0.9863523
```

```
cor(amazon, amazon_2)
```

```
##           [,1]
## AMZN.Adjusted 0.9727468
```

```
cor(amazon, amazon_3)
```

```
##           [,1]
## AMZN.Adjusted 0.9597422
```

```
google_1 = rep(0, (length(google)-1))
google_1[1:(length(google)-1)] = c(google[1:(length(google)-1)])
cor(x = as.vector(google[2:772]), y = google_1)
```

```
## [1] 0.9937932
```

```
cor(google[3:772], google[1:770])
```

```
##                GOOG.Adjusted
## GOOG.Adjusted    0.987505
```

```
cor(google[4:772], google[1:769])
```

```
##                GOOG.Adjusted
## GOOG.Adjusted    0.9823539
```

```
cor(microsoft[2:772], microsoft[1:771])
```

```
##                MSFT.Adjusted
## MSFT.Adjusted    0.9985696
```

```
cor(microsoft[3:772], microsoft[1:770])
```

```
##                MSFT.Adjusted
## MSFT.Adjusted    0.997512
```

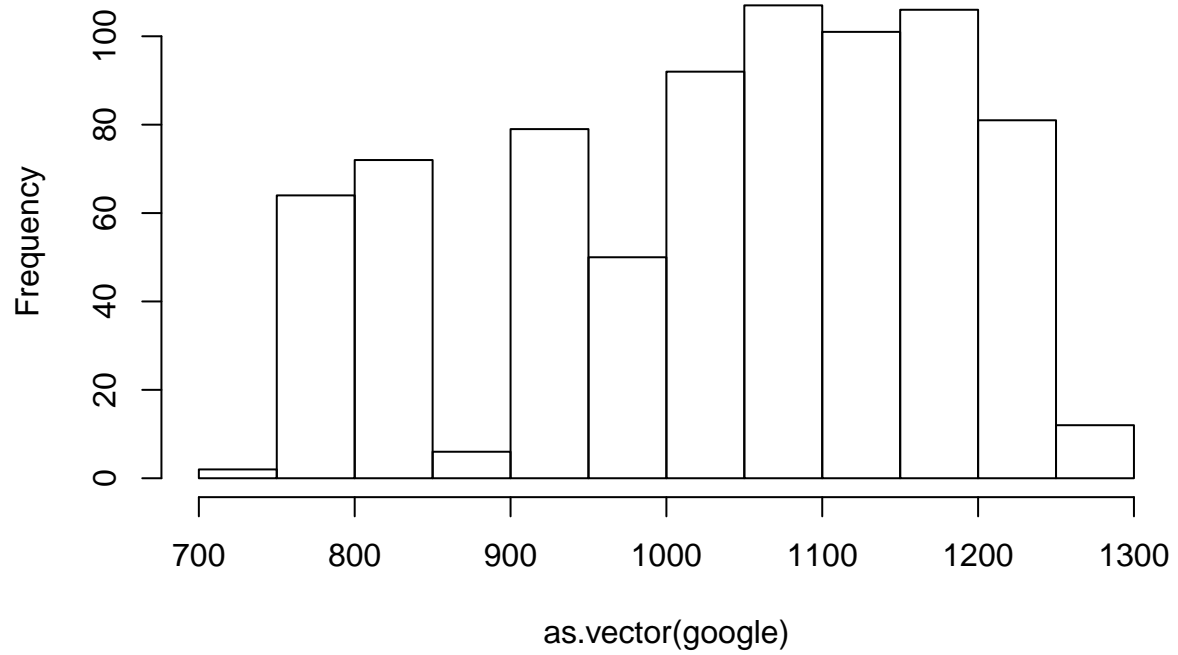
```
cor(microsoft[4:772], microsoft[1:769])
```

```
##                MSFT.Adjusted
## MSFT.Adjusted    0.9968224
```

All of them seem very positively autocorrelated. This was checked with 1-lag, 2-lag, and 3-lag.


```
hist(as.vector(google))
```

**Histogram of as.vector(google)**




```
hist(as.vector(microsoft))
```

**Histogram of as.vector(microsoft)**



```
hist(as.vector(amazon))
```





If the data had been stationary, we'd have seen consistent frequency histograms. Looking at the time series plot itself, we can easily infer that the mean is changing and they aren't stationary. Making a 50-50 split of the data will return different mean and variances.

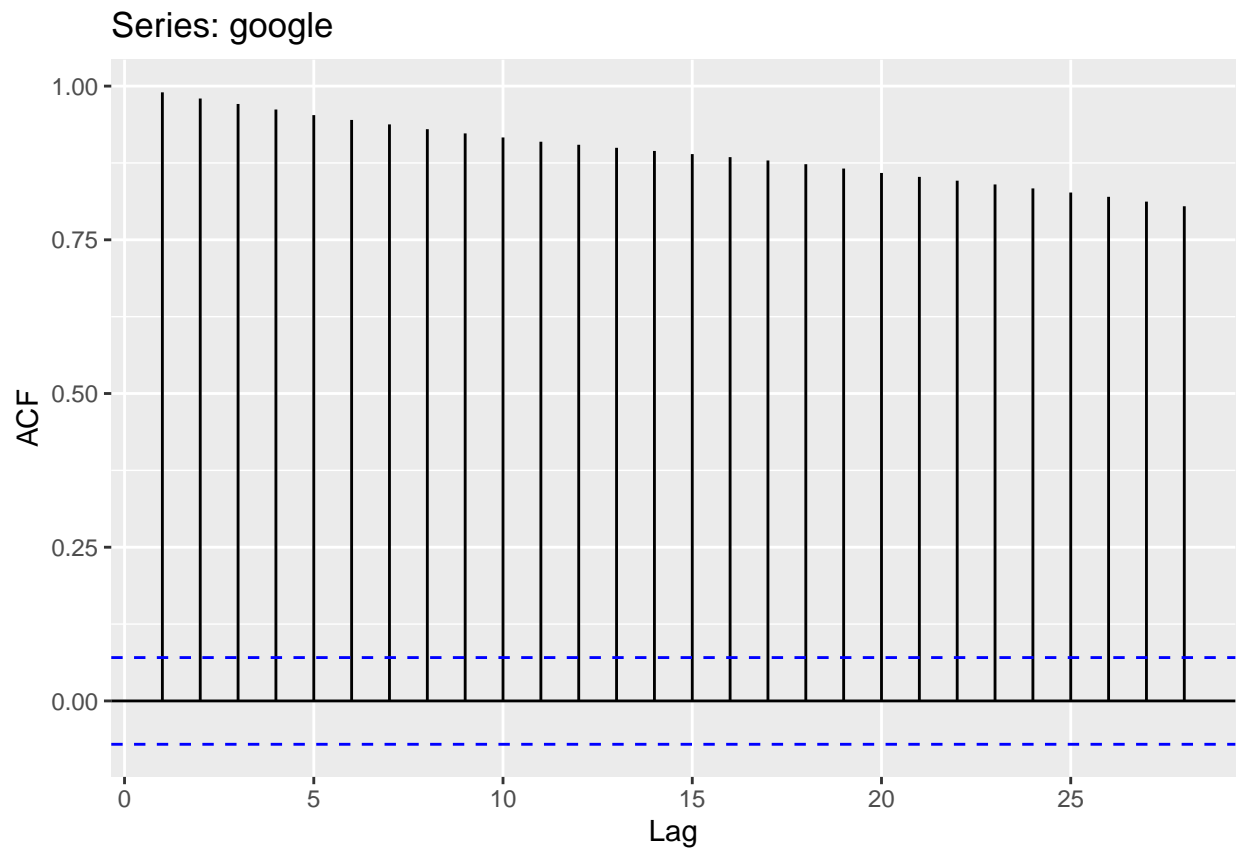
```
library(forecast)
```

```
## Registered S3 methods overwritten by 'forecast':
```

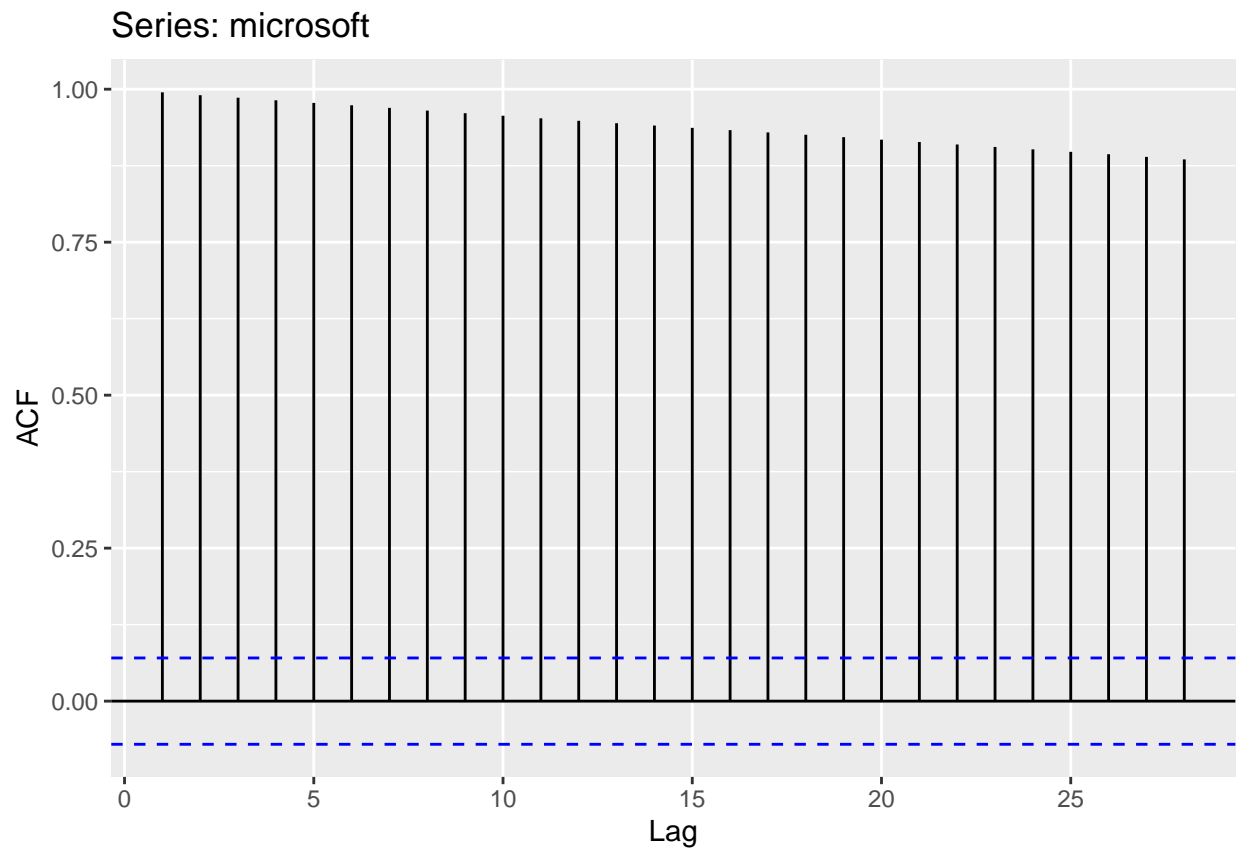
```
##   method          from  
##   fitted.fracdiff  fracdiff  
##   residuals.fracdiff fracdiff
```

```
library(tseries)
```

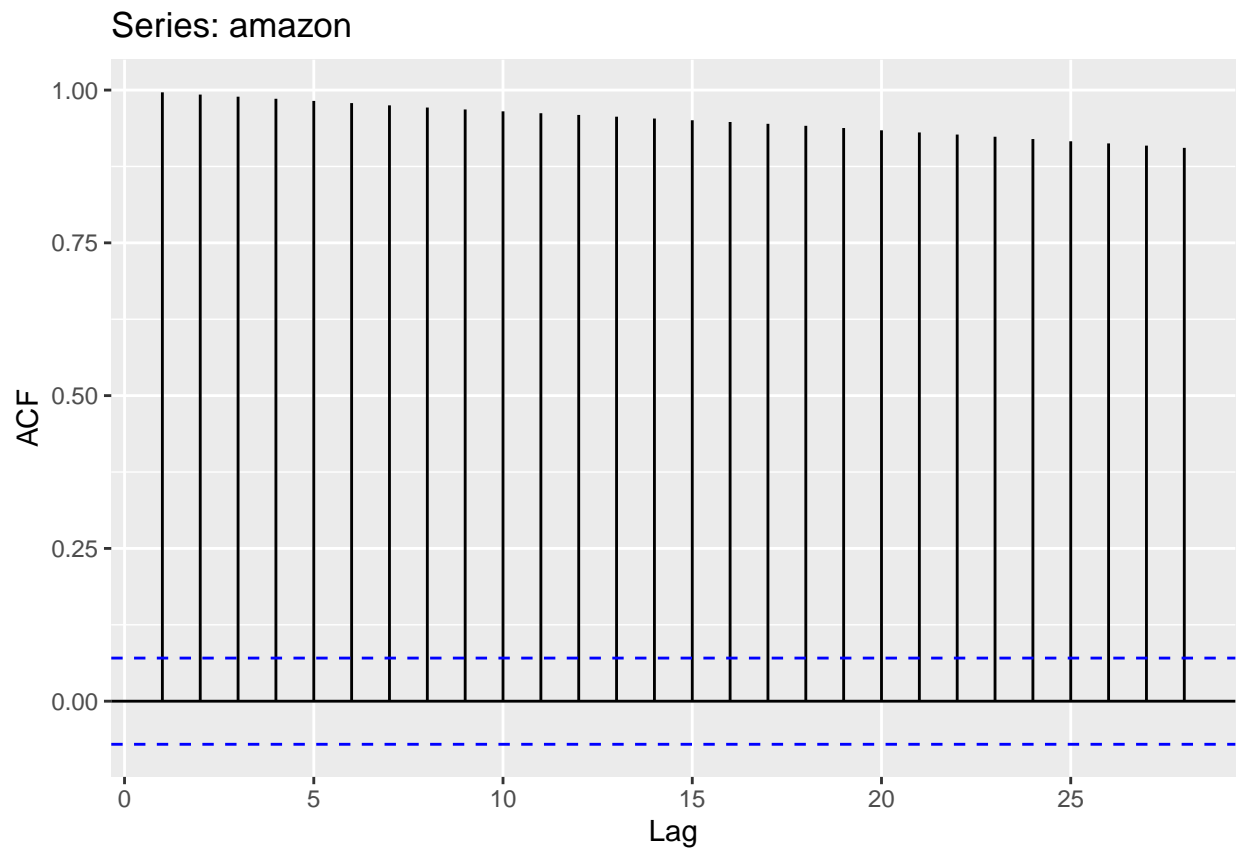
```
ggAcf(google)
```



```
ggAcf(microsoft)
```




```
ggAcf(amazon)
```



```
fcast = forecast(auto.arima(as.numeric(google)[1:617]), h = 155, level = c(95))  
plot(fcast)
```


## Forecasts from ARIMA(2,1,2)



It can be seen that the model isn't very good. We should smoothen the time series by the taking differences.


```
google_diff1 = rep(0, length(google)-1)
google_diff1 = as.vector(google)[2:772] - google[1:771]
fcast = forecast(auto.arima(as.numeric(google)[1:617]), h = 155, level = c(95))
plot(fcast)
```

## Forecasts from ARIMA(2,1,2)



```
fit_google = forecast(Arima(google[1:617], order = c(2,3,2)), h = 155, level = c(50))
plot(fit_google, ylim = c(700,2500), xlim = c(0,800))
par(new = T)
plot(as.vector(google), ylim = c(700,2500), xlim = c(0,800), xlab="", ylab="", main="", col="red", type
```


## Forecasts from ARIMA(2,3,2)



Both predictions on the Google stock have a high variability. The first fit seems better since it relates to the actual data much better. Doing the same analysis on different stocks.

```
fit_msft = forecast(Arima(microsoft[1:617], order = c(2,3,2)), h = 155, level = c(50))
plot(fit_msft, ylim = c(50, 160), xlim = c(0,800))
par(new = T)
plot(as.vector(microsoft), ylim = c(50,160), xlim = c(0,800), xlab="", ylab="", main="", col="red", typ
```

## Forecasts from ARIMA(2,3,2)

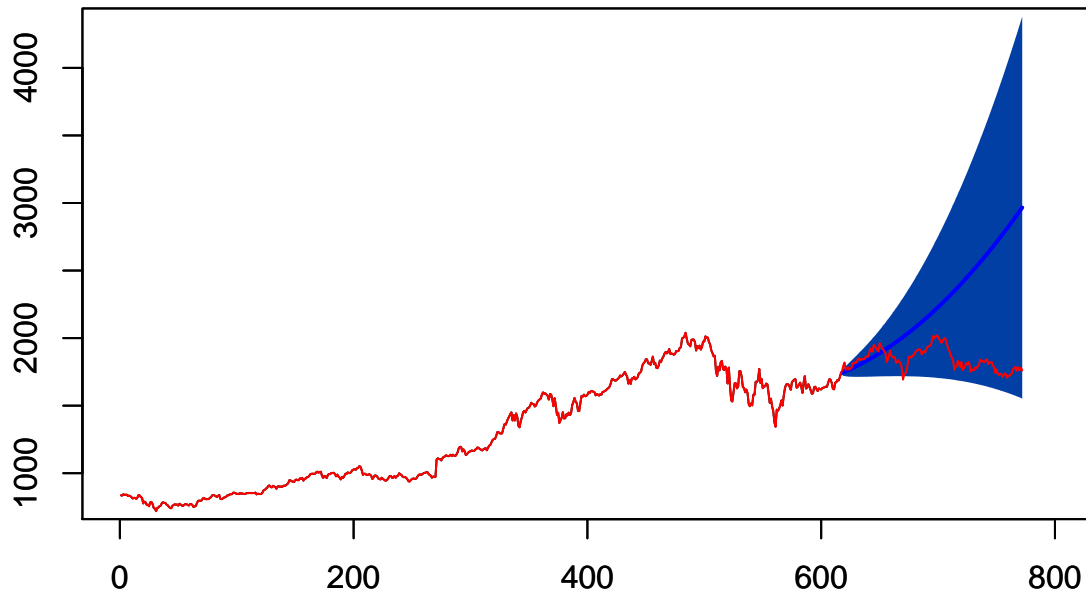


Microsoft's stock is predicted much better! Now for Amazon : –

```
fit_amzn = forecast(Arima(amazon[1:617], order = c(2,3,2)), h = 155, level = c(50))
plot(fit_amzn, xlim = c(0,800), ylim = c(800, 4300))
par(new = T)
plot(as.vector(amazon), ylim = c(800,4300), xlim = c(0,800), xlab="", ylab="", main="", col="red", type
```



## Forecasts from ARIMA(2,3,2)



Microsoft's stock was predicted the best with the least variance seen due to the confidence intervals. Amazon's was still not predicted very well when we check the actual data.

## Question 3

```
library(mlbench)
data("BreastCancer")
library(superml)
```

## Loading required package: R6

```
library(MASS)
```

Viewing the dataframe, we see that the Id column is a character column. All other are categorical variables encoded as numbers.

```
BreastCancer$Id = as.numeric(BreastCancer$Id)
lbl = superml::LabelEncoder$new()
BreastCancer$Class = lbl$fit_transform(BreastCancer$Class)
```

This encoding converts 'benign' to '0' and 'malignant' to '1' (retains the column as a factor).

Let's check for Spearman correlation with the dependent variables. We may very well discard the column *Id*.

```
cor(data.matrix(BreastCancer), method = 'spearman')
```

```
##               Id Cl.thickness  Cell.size  Cell.shape
## Id           1.000000000 -0.004408568 -0.04332721 -0.06009478
```

```
## Cl.thickness      -0.004408568  1.000000000  0.66646314  0.66412065
## Cell.size         -0.043327205  0.666463145  1.00000000  0.89190373
## Cell.shape        -0.060094775  0.664120648  0.89190373  1.00000000
## Marg.adhesion     -0.049512616  0.541592582  0.74251672  0.71173014
## Epith.c.size      -0.087263125  0.583978964  0.78715882  0.75923394
## Bare.nuclei       NA           NA           NA           NA
## Bl.cromatin        -0.095749854  0.538092726  0.71932965  0.69231583
## Normal.nucleoli   -0.070964365  0.570398358  0.75719871  0.72526781
## Mitoses           -0.075138482  0.418717438  0.50878003  0.47274889
## Class             -0.105498425  0.682451869  0.85548668  0.83639413
## Marg.adhesion     -0.04951262  -0.08726313  NA           -0.09574985
## Id                -0.04951262  -0.08726313  NA           -0.09574985
## Cl.thickness       0.54159258  0.58397896  NA           0.53809273
## Cell.size          0.74251672  0.78715882  NA           0.71932965
## Cell.shape         0.71173014  0.75923394  NA           0.69231583
## Marg.adhesion      1.00000000  0.66780890  NA           0.62451532
## Epith.c.size       0.66780890  1.00000000  NA           0.63950687
## Bare.nuclei        NA           NA           1           NA
## Bl.cromatin         0.62451532  0.63950687  NA           1.00000000
## Normal.nucleoli    0.63430951  0.70599721  NA           0.66230918
## Mitoses            0.44699240  0.48025536  NA           0.38694363
## Class              0.72799520  0.76273087  NA           0.74035037
## Normal.nucleoli    -0.07096437  -0.07513848  -0.1054984
## Id                -0.07096437  -0.07513848  -0.1054984
## Cl.thickness       0.57039836  0.41871744  0.6824519
## Cell.size          0.75719871  0.50878003  0.8554867
## Cell.shape         0.72526781  0.47274889  0.8363941
## Marg.adhesion      0.63430951  0.44699240  0.7279952
## Epith.c.size       0.70599721  0.48025536  0.7627309
## Bare.nuclei        NA           NA           NA
## Bl.cromatin         0.66230918  0.38694363  0.7403504
## Normal.nucleoli    1.00000000  0.50414036  0.7438226
## Mitoses            0.50414036  1.00000000  0.5267662
## Class              0.74382258  0.52676617  1.0000000
```

*Cell.size* and *Cell.shape* show very good correlations ( $\rho > 0.8$ ). *Bare.nuclei* contains some NAs. Let's remove them and check the correlation.

```
cor(data.matrix(BreastCancer), method = 'spearman')
```

```
## Id Cl.thickness Cell.size Cell.shape
## Id      1.000000000 -0.004408568 -0.04332721 -0.06009478
## Cl.thickness -0.004408568  1.000000000  0.66646314  0.66412065
## Cell.size   -0.043327205  0.666463145  1.00000000  0.89190373
## Cell.shape  -0.060094775  0.664120648  0.89190373  1.00000000
## Marg.adhesion -0.049512616  0.541592582  0.74251672  0.71173014
## Epith.c.size -0.087263125  0.583978964  0.78715882  0.75923394
## Bare.nuclei   NA           NA           NA           NA
## Bl.cromatin   -0.095749854  0.538092726  0.71932965  0.69231583
## Normal.nucleoli -0.070964365  0.570398358  0.75719871  0.72526781
## Mitoses       -0.075138482  0.418717438  0.50878003  0.47274889
## Class         -0.105498425  0.682451869  0.85548668  0.83639413
## Marg.adhesion -0.04951262  -0.08726313  NA           -0.09574985
## Id           -0.04951262  -0.08726313  NA           -0.09574985
## Cl.thickness  0.54159258  0.58397896  NA           0.53809273
```

```
## Cell.size      0.74251672  0.78715882      NA  0.71932965
## Cell.shape     0.71173014  0.75923394      NA  0.69231583
## Marg.adhesion  1.00000000  0.66780890      NA  0.62451532
## Epith.c.size   0.66780890  1.00000000      NA  0.63950687
## Bare.nuclei    NA          NA          1      NA
## Bl.cromatin    0.62451532  0.63950687      NA  1.00000000
## Normal.nucleoli 0.63430951  0.70599721      NA  0.66230918
## Mitoses        0.44699240  0.48025536      NA  0.38694363
## Class          0.72799520  0.76273087      NA  0.74035037
##               Normal.nucleoli  Mitoses  Class
## Id                -0.07096437 -0.07513848 -0.1054984
## Cl.thickness      0.57039836  0.41871744  0.6824519
## Cell.size         0.75719871  0.50878003  0.8554867
## Cell.shape        0.72526781  0.47274889  0.8363941
## Marg.adhesion     0.63430951  0.44699240  0.7279952
## Epith.c.size      0.70599721  0.48025536  0.7627309
## Bare.nuclei       NA          NA          NA
## Bl.cromatin        0.66230918  0.38694363  0.7403504
## Normal.nucleoli    1.00000000  0.50414036  0.7438226
## Mitoses            0.50414036  1.00000000  0.5267662
## Class              0.74382258  0.52676617  1.0000000
```

$\rho = 0.74$  for *bare.nuclei*.

Now, since the response variable is a binary variable, we'll use the binomial distribution as the link function.

```
for (i in 1:11){
  BreastCancer[,i] = as.numeric(BreastCancer[,i])
}
```

```
BreastCancer = na.omit(BreastCancer)
```

```
model = stats::glm(Class ~ . , data = BreastCancer, family = binomial(link = "logit"), control = list(
summary(model)
```

```
##
## Call:
## stats::glm(formula = Class ~ . , family = binomial(link = "logit"),
##   data = BreastCancer, control = list(maxit = 100))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.4892  -0.1155  -0.0613   0.0222   2.4672
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.016e+01  1.453e+00  -6.989 2.76e-12 ***
## Id           3.997e-08  7.336e-07   0.054 0.956546
## Cl.thickness  5.352e-01  1.419e-01   3.771 0.000163 ***
## Cell.size    -6.481e-03  2.092e-01  -0.031 0.975290
## Cell.shape    3.228e-01  2.308e-01   1.399 0.161836
## Marg.adhesion 3.307e-01  1.234e-01   2.679 0.007385 **
## Epith.c.size  9.640e-02  1.567e-01   0.615 0.538490
## Bare.nuclei   3.839e-01  9.548e-02   4.021 5.79e-05 ***
## Bl.cromatin   4.479e-01  1.716e-01   2.609 0.009070 **
## Normal.nucleoli 2.135e-01  1.131e-01   1.887 0.059221 .
```

```
## Mitoses          5.382e-01  3.263e-01  1.650 0.099035 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 884.35  on 682  degrees of freedom
## Residual deviance: 102.90  on 672  degrees of freedom
## AIC: 124.9
##
## Number of Fisher Scoring iterations: 8
```

The statistics displayed by the `stats::glm` function displays : –

1. **Call** - What actual command was given to generate the generalised linear model.
2. **Deviance Residuals** - Gives the summary of the deviance residuals - their minimum value, 25th, 50th, and 75th percentile, and the maximum value.
3. **Coefficients** - These are the coefficients (or slopes) given to the different independent variables. Note the significance codes below. Our data has the intercept, Cl.thickness, Marg.adhesion, Bare.nuclei, Bl.cromatin as significant variables. These coefficients are actually in units of logit. These are log of odd ratios.
4. **Null and Residual deviance** - The null model (containing random values) had a deviance of 884.35 on 682 degrees of freedom. The residual deviance could bring it down to 102.90 — a significant drop with a reduction to 672 degrees of freedom. Lesser this deviance, better the fit.
5. **\*AIC** - The Akaike Information Criteria or AIC is 124.9. This model can be compared with other such models on the basis of complexity with this measure.
6. **Number of Fisher Scoring Iterations** - The maximum iterations were given to be 100. This number tells us that the model was itself able to converge (using Newton's approximation method) with 8 iterations.

```
exp(cbind(coef(model), confint(model, level = 0.95)))
```

```
## Waiting for profiling to be done...
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
##              2.5 %      97.5 %
## (Intercept)  3.883384e-05 2.026971e-06 0.0006751557
## Id          1.000000e+00 9.999981e-01 1.0000006406
## Cl.thickness 1.707765e+00 1.316023e+00 2.3112071985
## Cell.size    9.935402e-01 6.734836e-01 1.5505839408
## Cell.shape   1.381011e+00 8.610222e-01 2.1567226486
## Marg.adhesion 1.391920e+00 1.097572e+00 1.7984950474
## Epith.c.size  1.101202e+00 8.047423e-01 1.4992552531
## Bare.nuclei   1.468056e+00 1.226733e+00 1.7906219676
## Bl.cromatin   1.564976e+00 1.131276e+00 2.2279653149
## Normal.nucleoli 1.237953e+00 9.981035e-01 1.5630323124
## Mitoses       1.712889e+00 9.914547e-01 3.0297864371
```

The warnings displayed is due to the fact of ‘perfect distinguishability’, i.e., the data is very able to separate the benign and malignant cases. This gives us exact 0 or 1 probabilities in two cases. The odds discussed earlier were in logit units or log odd ratios, so we exponentiate them. They give us 95% confidence intervals (2.5% on both sides). Here we see that all the variables are positively correlated with our target variable - meaning as they increase from the discrete 0 towards 10, the target variable increases from 0 to 1 (benign to malignant). The biggest odd ratios are seen in *Cell.size* and the intercept. This means that *Cell.size*

dominates the prediction towards benign or malignancy. All other have a near about coefficient of 1 (although the confidence interval seems too high) so they are being used as additives in the model.